When Do Answers Change? Estimating Question Recency Demands in QA

Peter Schulze
https://github.com/peterschulze04
Peter.Schulze@student.uibk.ac.at

Fabian Stiewe
https://github.com/Fabianstw
Fabian.Stiewe@student.uibk.ac.at

Abstract

Large language models (LLMs) answer questions using knowledge that may change over time; however, most evaluations remain static. This work explores the demands of question recency, specifically the frequency with which answers must be updated to maintain accuracy, and the precision with which models can predict the necessary update frequency. We construct a *generation-and-labeling* pipeline that produces recency-aware questions along two axes: stationary vs.versus non-stationary temporal behaviour, and single-event vs. multi-event dependencies. The multi-event category captures questions whose answers rely on multiple temporally connected events, including causal and purely co-occurring pairs. The single-event category focuses on questions about a single evolving process. Stationarity labels distinguish predictable update cycles from volatile changes. Employing a Llama model, the pipeline generates and labels a 600-question dataset, assigning recency classes, stationarity labels, and multi-event markers. Subsequently, it applies perturbations to construct more challenging temporal variants. The resulting dataset extends prior work, providing new and more demanding questions with which to evaluate LLMs on when those answers change.

1 Introduction

In the last couple of years, large language models (LLMs) have demonstrated remarkable capabilities across a wide range of natural language processing tasks. In particular, their performance on question answering (QA) tasks has been studied extensively [Fischer et al., 2024, Brown et al., 2020], with very strong results. However, a persistent challenge for QA systems is handling temporal dynamics, specifically answering questions whose correct answers change over time.

Much of this prior work largely ignores the fact that many real-world questions require information that is time-sensitive and may change frequently, such as:

- *"What is the current population of Germany?"*
- *"Who is the current president of the United States?"*

As facts evolve, answers that were once correct inevitably degrade, and outdated responses can harm user trust. Despite this, many QA evaluations remain static, implicitly assuming that factual correctness is invariant over time.

Prior work on temporal QA and time-aware modeling has highlighted this limitation from several perspectives. These include datasets for time-sensitive questions [Chen et al., 2021], temporal knowledge bases for language models [Dhingra et al., 2022], real-time QA benchmarks [Kasai et al., 2024], search-augmented approaches for refreshing model knowledge [Vu et al., 2024], and self-updating present-anchored QA suites [Meem et al., 2024]. While these efforts underscore the importance of temporal awareness, they typically treat time sensitivity as a coarse or binary property and do not explicitly characterize how frequently answers change or how complex their temporal dependencies are.

We refer to this requirement as the *recency demand* of a question: the degree to which an answer must be updated over time to remain accurate. Recency demand varies across questions, depending on factors such as the predictability of change and whether an answer depends on a single evolving event or on multiple temporally related events.

A very recent paper, *RecencyQA* (unpublished), introduces a dataset explicitly labeled with recency demands. The authors annotate questions according to how frequently their answers are expected to change and use this dataset to evaluate the ability of various LLMs to predict a question's recency demand. Their results indicate that, while models capture coarse distinctions between static and

Dataset	Creation	KC	Multi-hop	Recency-Label	Multi-Event	# Ques.
TimeQA [Chen et al., 2021]	Templ.-Wikidata	Wikipedia	✗	✗	✗	20 000
SituatedQA [Zhang and Choi, 2021]	Man.-Filt.	Wikipedia	✗	✗	✗	12 000
TempLama [Dhingra et al., 2022]	Templ./Cloze	Custom-News	✗	✗	✗	50 000
StreamingQA [Liška et al., 2022]	Man.+Gen.	WMT News	✓	✗	✗	410 000
ArchivalQA [Wang et al., 2022]	Gen.	NYT Articles	✗	✗	✗	532 000
ChroniclingAmericaQA [Piryani et al., 2024]	Gen.	Chronicling America Newspapers	✗	✗	✗	485 000
RealTimeQA [Kasai et al., 2024]	News websites	News Articles	✓	✗	✗	∼5000
PATQA [Meem et al., 2024]	Templ.-wikidata	Wikipedia	✓	✗	✗	6172
FreshQA [Vu et al., 2024]	Man.	Google Search	✓	✗	✗	600
RecencyQA (*unpublished*)	Man.-Filt.+Gen	Wikipedia/Wikidata	✓	✓	✗	6115
RecencyQA-Multi (*ours*)	Man.-Filt.+Gen	RecencyQA+LLM	✓	✓	✓	1432

Table 1: Overview of related datasets.

rapidly changing facts, accurately estimating finer-grained recency requirements remains challenging.

Building on this line of work, our approach extends prior recency-aware QA datasets along several important dimensions. Rather than relying on a fixed set of annotated questions, we introduce a generation-and-labeling pipeline that systematically expands a small seed set into a substantially richer collection of temporal questions. Starting from 75 input questions, the pipeline produces approximately 1500 recency-aware questions by varying both temporal behavior and event structure. In addition to distinguishing stationary from non-stationary temporal dynamics, we explicitly control for event dependency—an aspect that has received little explicit attention in prior work—by generating single-event questions as well as multi-event questions involving two or three temporally connected events, with both causal and non-causal relationships. Each generated question is further annotated with one or more recency labels and contextual conditions under which the label applies. This structured expansion enables finer-grained analysis of temporal complexity in QA than prior datasets, which typically focus on isolated questions or coarse recency distinctions.

We start explaining the generation pipeline in Section 2, followed by testing various LLMs on the created dataset in Section 3. Using the results from the experiment, we fine-tune a smaller model and compare it to the larger LLMs in Section 4. Finally, we conclude with a discussion of findings and future work in Section 5.

2 Dataset Creation

Our dataset pipeline converts a small seed set of temporal questions into a larger collection annotated with fine-grained recency demands and explicit event cardinality. We implement the pipeline with Llama-3.3-70B-Instruct-Turbo [Touvron et al., 2023] via Together AI[1], using deterministic prompts that separate generation and labeling (full templates appear in Appendix A). This separation keeps the phrasing space broad while making the labeling step auditable and repeatable.

2.1 Seed data and model setup

We start from a compact seed split of RecencyQA. Each seed question is passed as a one-shot example to the generator to anchor topic diversity and phrasing. The generation and labeling calls use the Together chat API with temperature 1.0 and top-p 0.95. Token usage is logged to track cost, but no seed answers are exposed to the model.

2.2 Question taxonomy

We factor questions along stationarity and event cardinality, storing both the categorical choices and the cardinality as explicit metadata for every record.

- **Stationary vs. non-stationary**: Stationary questions have predictable update cycles; non-stationary questions depend on volatile or event-triggered changes.
- **Event cardinality (1/2/3 events)**: Single-event questions revolve around one evolving process. Multi-event questions use either two or three events. For multi-event items we prompt both *causal* chains and *temporal-only* co-occurrences across different domains. Each record stores an event dependency flag (Single-Event vs. Multi-Event), the number of events, and the generation type when multi-event.

This yields ten prompt modes: single-event stationary and non-stationary, two-event causal and

[1]Together AI | The AI Native Cloud: https://www.together.ai/

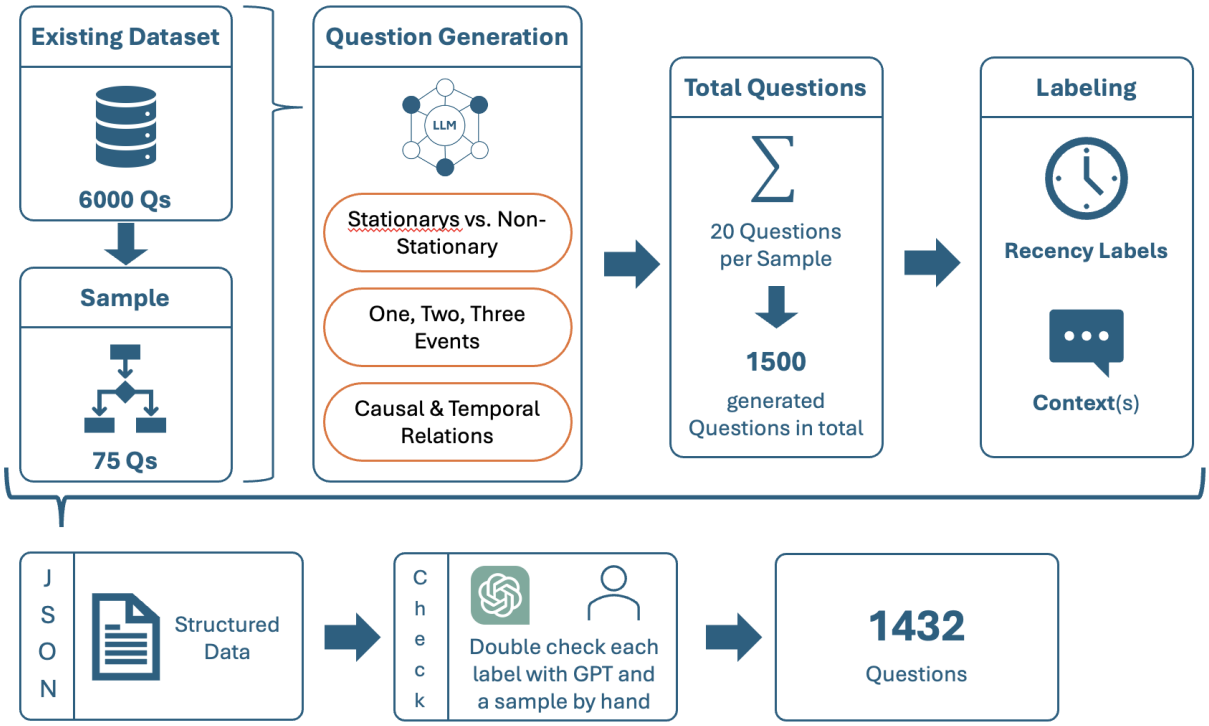

Figure 1: Caption text.

temporal-only counterparts for each stationarity setting, and three-event causal and temporal-only counterparts. The metadata reflects these choices through an event dependency flag (Single-Event vs. Multi-Event), the number of events, and for multi-event items a generation type tag (causal vs. temporal-only).

2.3 Question generation

Prompt templates instruct the model to output exactly two questions per call and forbid placeholders or paraphrasing of the seed (Appendix A.1). A helper repeats calls until the target count is reached for each mode, reissuing only the missing portion to avoid over-generation. The ten prompt families are: stationary single (Appendix A.1.1); stationary two-event causal (Appendix A.1.2) and temporal-only (Appendix A.1.3); stationary three-event causal (Appendix A.1.4) and temporal-only (Appendix A.1.5); non-stationary single (Appendix A.1.6); non-stationary two-event causal (Appendix A.1.7) and temporal-only (Appendix A.1.8); and non-stationary three-event causal (Appendix A.1.9) and temporal-only (Appendix A.1.10). Every generated question is tagged with the seed identifier and metadata for stationarity, event dependency, number of events, and (when applicable) whether the multi-event relation was prompted as causal or temporal-only.

2.4 Recency labeling

Each question is labeled with recency classes that describe how quickly its answer should be refreshed. We use the following twelve labels:

An hour	A few hours
A day	A few days
A week	A few weeks
A month	A few months
A year	A few years
Many years	Never

Stationary questions receive one label; non-stationary questions receive two labels to capture alternative temporal regimes. For every label we also elicit a one-sentence temporal context (event, phase, or condition) that makes the question relevant. The labeling prompt prohibits meta-reasoning or explicit time stamps and is reproduced in Appendix A.2 via the stationary template (Appendix A.2.1) and the non-stationary template (Appendix A.2.2). Labels and contexts are merged into a structured list so downstream models can pair each recency class with its motivating scenario. We retain only entries whose labels parse as JSON and whose context count matches the expected label count.

2.5 Sampling and size

Running the ten-generation-mode configuration on 75 seed questions yields 1,498 retained questions after discarding malformed generations that fail JSON parsing or lack labels. The ensure-and-top-off loop keeps modes balanced: 748 stationary vs. 750 non-stationary questions; 300 single-event, 598 two-event, and 600 three-event items; and within multi-event items 600 were prompted as causal chains and 598 as temporal-only pairs or triples. Mode totals differ from the ideal 1,500 by two items that were filtered for malformed JSON.

2.6 Output format

All examples are written as one JSON object per line. Each record stores the seed question identifier, question text, stationarity flag, event dependency type, number of events, the causal vs. temporal-only generation tag for multi-event items, and the list of recency labels with their paired contexts. The resulting corpus offers controlled variation across temporal behaviors and provides explicit supervision for estimating when an answer is likely to drift. We also keep the original seed linkage so downstream work can group generated questions by provenance or remove seed-specific leakage.

2.7 Dataset Statistics

The released split contains 1,498 examples. Stationarity is balanced (748 stationary, 750 non-stationary), and label cardinality follows the prompt contract (one label for stationary, two for non-stationary). Event cardinality is spread across 300 single-event, 598 two-event, and 600 three-event questions, with multi-event items split nearly evenly between causal (600) and temporal-only (598) prompting. The most common recency labels are "A-Year" (406), "A-Day" (368), and "A-Few-Days" (341); the tail includes a handful of model-emitted variants (e.g., "A-Quarter", "A-Few-Minutes") that remain in the data for transparency and can be remapped if a downstream system requires a closed label set.

3 Experiments

4 LLM Fine-Tuning

5 Conclusion and Future Work

References

Tom B. Brown, Benjamin Mann, Nick Ryder, Melanie Subbiah, Jared Kaplan, Prafulla Dhariwal, Arvind Neelakantan, Pranav Shyam, Girish Sastry, Amanda Askell, Sandhini Agarwal, Ariel Herbert-Voss, Gretchen Krueger, Tom Henighan, Rewon Child, Aditya Ramesh, Daniel M. Ziegler, Jeffrey Wu, Clemens Winter, Christopher Hesse, Mark Chen, Eric Sigler, Mateusz Litwin, Scott Gray, Benjamin Chess, Jack Clark, Christopher Berner, Sam McCandlish, Alec Radford, Ilya Sutskever, and Dario Amodei. Language models are few-shot learners, 2020. URL https://arxiv.org/abs/2005.14165.

Wenhu Chen, Xinyi Wang, and William Yang Wang. A dataset for answering time-sensitive questions, 2021. URL https://arxiv.org/abs/2108.06314.

Bhuwan Dhingra, Jeremy R. Cole, Julian Martin Eisenschlos, Daniel Gillick, Jacob Eisenstein, and William W. Cohen. Time-aware language models as temporal knowledge bases. *Transactions of the Association for Computational Linguistics*, 10:257–273, 03 2022. ISSN 2307-387X. doi: 10.1162/tacl_a_00459. URL https://doi.org/10.1162/tacl_a_00459.

Kevin Fischer, Darren Fürst, Sebastian Steindl, Jakob Lindner, and Ulrich Schäfer. Question: How do large language models perform on the question answering tasks? answer:, 2024. URL https://arxiv.org/abs/2412.12893.

Jungo Kasai, Keisuke Sakaguchi, Yoichi Takahashi, Ronan Le Bras, Akari Asai, Xinyan Yu, Dragomir Radev, Noah A. Smith, Yejin Choi, and Kentaro Inui. Realtime qa: What's the answer right now?, 2024. URL https://arxiv.org/abs/2207.13332.

Adam Liška, Tomáš Kočiský, Elena Gribovskaya, Tayfun Terzi, Eren Sezener, Devang Agrawal, Cyprien de Masson d'Autume, Tim Scholtes, Manzil Zaheer, Susannah Young, Ellen Gilsenan-McMahon, Sophia Austin, Phil Blunsom, and Angeliki Lazaridou. Streamingqa: A benchmark for adaptation to new knowledge over time in question answering models, 2022. URL https://arxiv.org/abs/2205.11388.

Jannat Meem, Muhammad Rashid, Yue Dong, and Vagelis Hristidis. PAT-questions: A self-updating benchmark for present-anchored temporal question-answering. In Lun-Wei Ku, Andre Martins, and Vivek Srikumar, editors, *Findings of the Association for Computational Linguistics: ACL 2024*, pages 13129–13148, Bangkok, Thailand, August 2024. Association for Computational Linguistics. doi: 10.18653/v1/2024.findings-acl.777. URL https://aclanthology.org/2024.findings-acl.777/.

Bhawna Piryani, Jamshid Mozafari, and Adam Jatowt. Chroniclingamericaqa: A large-scale question answering dataset based on historical american newspaper pages. In *Proceedings of the 47th International ACM SIGIR Conference on Research and Development in Information Retrieval*, SIGIR 2024, page 2038–2048. ACM, July 2024. doi: 10.1145/3626772.3657891. URL http://dx.doi.org/10.1145/3626772.3657891.

Hugo Touvron, Thibaut Lavril, Gautier Izacard, Xavier Martinet, Marie-Anne Lachaux, Timothée Lacroix, Baptiste Rozière, Naman Goyal, Eric Hambro, Faisal Azhar, Aurelien Rodriguez, Armand Joulin, Edouard Grave, and Guillaume Lample. Llama: Open and efficient foundation language models, 2023. URL https://arxiv.org/abs/2302.13971.

Tu Vu, Mohit Iyyer, Xuezhi Wang, Noah Constant, Jerry Wei, Jason Wei, Chris Tar, Yun-Hsuan Sung, Denny Zhou, Quoc Le, and Thang Luong. FreshLLMs: Refreshing large language models with search engine augmentation. In Lun-Wei Ku, Andre Martins, and Vivek Srikumar, editors, *Findings of the Association for Computational Linguistics: ACL 2024*, pages 13697–13720, Bangkok, Thailand, August 2024. Association for Computational Linguistics. doi: 10.18653/v1/2024.findings-acl.813. URL https://aclanthology.org/2024.findings-acl.813/.

Jiexin Wang, Adam Jatowt, and Masatoshi Yoshikawa. Archivalqa: A large-scale benchmark dataset for open domain question answering over historical news collections, 2022. URL https://arxiv.org/abs/2109.03438.

Michael J. Q. Zhang and Eunsol Choi. Situatedqa: Incorporating extra-linguistic contexts into qa, 2021. URL https://arxiv.org/abs/2109.06157.

A Prompt Templates

A.1 Generation prompts

A.1.1 Stationary single-event

```
You generate STATIONARY temporal questions.

Definition (internal):
- A stationary question has a stable temporal update behavior.
- The answer changes over time, but the timespan how often the answer must be updated
  would remain the same regardless of when the question is asked.

Task:
Generate EXACTLY 2 stationary temporal questions focusing on ONE event or process.

Rules:
- Must rely on real-world phenomena that change over time.
- Must be stable, cyclical, predictable, or rhythm-based.
- Must NOT paraphrase the examples.
- Must NOT use placeholders.
- Do NOT mention stationarity in the question.

Examples:
{examples}

Return JSON:
{
  "questions": ["q1","q2"]
}
```

A.1.2 Stationary multi-event (causal)

```
You generate STATIONARY multi-event temporal questions.

Task:
Generate EXACTLY 2 stationary temporal questions involving TWO events that are causally related.

Rules:
- Must involve at least TWO distinct temporal events.
- Events must have a clear cause-effect relationship.
- Temporal behavior must be stable, predictable, cyclical, or regular.
- Must rely on real-world change.
- Do NOT paraphrase examples.
- Do NOT use placeholders.
- Do NOT mention stationarity.

Examples:
{examples}

Return JSON:
{
  "questions": ["q1","q2"]
}
```

A.1.3 Stationary multi-event (temporal-only)

```
You generate STATIONARY multi-event temporal questions.

Task:
Generate EXACTLY 2 stationary temporal questions involving TWO events that are ONLY temporally
↪  connected.

Rules:
- Events must be from clearly different real-world domains.
- Events must NOT influence each other causally.
- Must NOT belong to the same topic, organization, or event series.
- Temporal behavior must be stable, predictable, cyclical, or regular.
- Must rely on real-world change.
- Do NOT paraphrase examples.
- Do NOT use placeholders.
- Do NOT mention stationarity.
```

```
Examples:
{examples}

Return JSON:
{
  "questions": ["q1","q2"]
}
```

A.1.4 Stationary three-event (causal)

```
You generate STATIONARY multi-event temporal questions.

Task:
Generate EXACTLY 2 stationary temporal questions involving THREE events
that are causally related in a chain or network.

Rules:
- Must involve EXACTLY THREE distinct temporal events.
- Events must have clear cause-effect relationships.
- Temporal behavior must be stable, predictable, cyclical, or regular.
- Must rely on real-world change.
- Do NOT paraphrase examples.
- Do NOT use placeholders.
- Do NOT mention stationarity.

Examples:
{examples}

Return JSON:
{
  "questions": ["q1", "q2"]
}
```

A.1.5 Stationary three-event (temporal-only)

```
You generate STATIONARY multi-event temporal questions.

Task:
Generate EXACTLY 2 stationary temporal questions involving THREE events
that are ONLY temporally connected.

Rules:
- Must involve EXACTLY THREE distinct events.
- Events must be from clearly different real-world domains.
- Events must NOT influence each other causally.
- Must NOT belong to the same topic, organization, or event series.
- Temporal behavior must be stable, predictable, cyclical, or regular.
- Must rely on real-world change.
- Do NOT paraphrase examples.
- Do NOT use placeholders.
- Do NOT mention stationarity.

Examples:
{examples}

Return JSON:
{
  "questions": ["q1", "q2"]
}
```

A.1.6 Non-stationary single-event

```
You generate NON-STATIONARY temporal questions.

Definition (internal):
- A non-stationary question has unstable temporal update behavior.
- How frequently the answer must be updated depends strongly on WHEN the question is asked.
- OR the question is only relevant within short, event-dependent windows.

Task:
Generate EXACTLY 2 non-stationary temporal questions focusing on ONE event.
```

```
Rules:
- Must rely on quickly evolving or unstable processes.
- Must be meaningful and real-world.
- Do NOT paraphrase examples.
- Do NOT use placeholders.
- Do NOT mention non-stationarity explicitly.

Examples:
{examples}

Return JSON:
{
  "questions": ["q1","q2"]
}
```

A.1.7 Non-stationary multi-event (causal)

```
You generate NON-STATIONARY multi-event temporal questions.

Task:
Generate EXACTLY 2 non-stationary temporal questions involving TWO events that are causally related.

Rules:
- At least one event must be unstable, unpredictable, or highly time-sensitive.
- Events must have a clear cause-effect relationship.
- Must rely on real-world temporal change.
- Must involve at least TWO distinct temporal events.
- Do NOT paraphrase examples.
- Do NOT use placeholders.
- Do NOT mention non-stationarity.

Examples:
{examples}

Return JSON:
{
  "questions": ["q1","q2"]
}
```

A.1.8 Non-stationary multi-event (temporal-only)

```
You generate NON-STATIONARY multi-event temporal questions.

Task:
Generate EXACTLY 2 non-stationary temporal questions involving TWO events that are ONLY temporally
↪  connected.

Rules:
- Events must be from clearly different real-world domains.
- Events must NOT influence each other causally.
- Must NOT belong to the same topic, organization, or event series.
- At least one event must be unstable, unpredictable, or highly time-sensitive.
- Must rely on real-world temporal change.
- Must involve at least TWO distinct temporal events.
- Do NOT paraphrase examples.
- Do NOT use placeholders.
- Do NOT mention non-stationarity.

Examples:
{examples}

Return JSON:
{
  "questions": ["q1","q2"]
}
```

A.1.9 Non-stationary three-event (causal)

```
You generate NON-STATIONARY multi-event temporal questions.

Task:
```

```
Generate EXACTLY 2 non-stationary temporal questions involving THREE events
that are causally related.

Rules:
- Must involve EXACTLY THREE distinct temporal events.
- At least one event must be unstable, unpredictable, or highly time-sensitive.
- Events must have clear cause-effect relationships.
- Must rely on real-world temporal change.
- Do NOT paraphrase examples.
- Do NOT use placeholders.
- Do NOT mention non-stationarity.

Examples:
{examples}

Return JSON:
{
  "questions": ["q1", "q2"]
}
```

A.1.10 Non-stationary three-event (temporal-only)

```
You generate NON-STATIONARY multi-event temporal questions.

Task:
Generate EXACTLY 2 non-stationary temporal questions involving THREE events
that are ONLY temporally connected.

Rules:
- Must involve EXACTLY THREE distinct temporal events.
- Events must be from clearly different real-world domains.
- Must NOT influence each other causally.
- At least one event must be unstable, unpredictable, or highly time-sensitive.
- Must rely on real-world temporal change.
- Do NOT paraphrase examples.
- Do NOT use placeholders.
- Do NOT mention non-stationarity.

Examples:
{examples}

Return JSON:
{
  "questions": ["q1", "q2"]
}
```

A.2 Labeling prompts

A.2.1 Stationary labeling

```
Analyze this temporal question and produce temporal labels:

"{question}"

Your tasks (internal reasoning only, output JSON only):

1. Provide ONE recency label:
{
  "An-Hour": "changes within one hour",
  "A-Few-Hours": "changes within a few hours",
  "A-Day": "changes within one day",
  "A-Few-Days": "changes within a few days",
  "A-Week": "changes within one week",
  "A-Few-Weeks": "changes within a few weeks",
  "A-Month": "changes within one month",
  "A-Few-Months": "changes within a few months",
  "A-Year": "changes within one year",
  "A-Few-Years": "changes within a few years",
  "Many-Years": "changes after 10 or more years",
  "Never": "never changes"
}
```

```
2. Based on the selected label, write a short temporal context (ONE sentence) describing the current
↪  event, phase, or condition in which the question is asked. This does not have to be related to the
↪  question causally.
   - Must describe an EVENT, PHASE or CONDITION, taking place when the question becomes relevant.
   - No specific years.
   - No meta reasoning.
   - No "current"
   - No "the question is asked"

Return JSON:
{
  "recency_list": ["label1"],
  "context_list": ["context1"]
}
```

A.2.2 Non-stationary labeling

```
Analyze this temporal question and produce temporal labels:

"{question}"

Your tasks (internal reasoning only, output JSON only):

1. Provide TWO recency labels (for two different realistic temporal situations), choose from:
{
  "An-Hour": "changes within one hour",
  "A-Few-Hours": "changes within a few hours",
  "A-Day": "changes within one day",
  "A-Few-Days": "changes within a few days",
  "A-Week": "changes within one week",
  "A-Few-Weeks": "changes within a few weeks",
  "A-Month": "changes within one month",
  "A-Few-Months": "changes within a few months",
  "A-Year": "changes within one year",
  "A-Few-Years": "changes within a few years",
  "Many-Years": "changes after 10 or more years",
  "Never": "never changes"
}

2. For each selected label provide ONE short temporal context (ONE sentence each) describing the
↪  current event, phase, or condition in which the question is asked. This does not have to be
↪  related to the question causally.
  - Each context must describe a different EVENT, PHASE or CONDITION, taking place when the question
  ↪  becomes relevant.
   - No specific years.
   - No meta reasoning.
   - No "current"
   - No "the question is asked"

Return JSON:
{
  "recency_list": ["label1", "label2"],
  "context_list": ["context1", "context2"]
}
```